

Finishing Up Instruction Formats & CALL (Compiler/Assembler/Linker/Loader)

U-Format for “Upper Immediate” instructions

31	12 11	7 6	0
imm[31:12]	rd	opcode	
20	5	7	
U-immediate[31:12]	dest	LUI	
U-immediate[31:12]	dest	AUIPC	

- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
 - LUI – Load Upper Immediate
 - AUIPC – Add Upper Immediate to PC

LUI to create long immediates

- LUI writes the upper 20 bits of the destination with the immediate value, and clears the lower 12 bits.
- Together with an ADDI to set low 12 bits, can create any 32-bit value in a register using two instructions (LUI/ADDI).
- Assembler uses this to create pseudo instructions

LUI x10, 0x87654 # x10 = 0x87654000

ADDI x10, x10, 0x321 # x10 = 0x87654321

One Corner Case

How to set 0xDEADBEEF?

```
LUI x10, 0xDEADB      # x10 = 0xDEADB000  
ADDI x10, x10, 0xEEF # x10 = 0xDEADAEEF
```

ADDI 12-bit immediate is *always* sign-extended, if top bit is set, will add a -1 from the upper 20 bits

Solution

How to set 0xDEADBEEF?

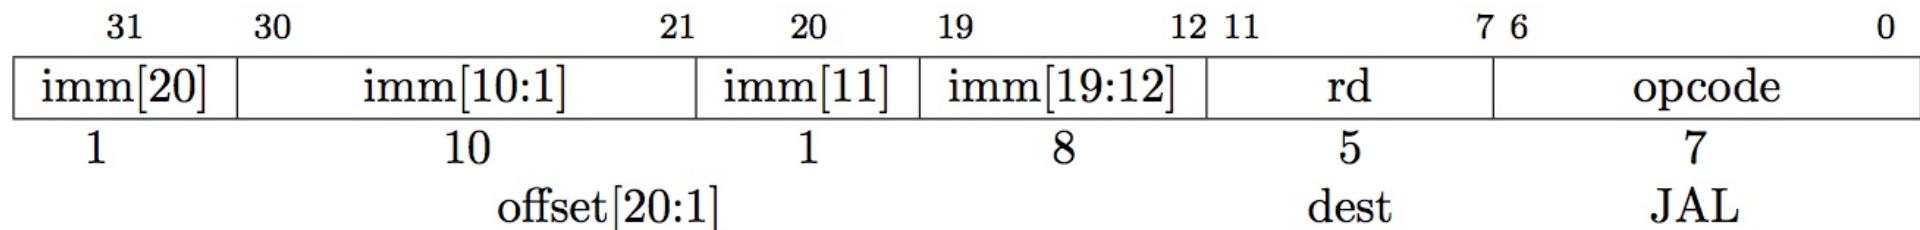
```
LUI x10, 0xDEADC      # x10 = 0xDEADC000  
ADDI x10, x10, 0xEEF # x10 = 0xDEADBEEF
```

Pre-increment value placed in upper 20 bits, if sign bit will be set on immediate in lower 12 bits.

Assembler pseudo-op handles all of this:

```
li x10, 0xDEADBEEF # Creates two instructions
```

J-Format for Jump Instructions



- JAL saves PC+4 in register rd (the return address)
 - Assembler “j” jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

Why is the offset that way?

- The j instruction encodes imm[20:1]... It fixes imm[0] as 0
 - So why is it always adding 2 and not 4?
I thought all instructions were 4 bytes and had to be word aligned...
 - So why not have it be imm[21:2] and the lower 2 bits 0?
- Well, they aren't for RISC-V (even if they are for us)
 - RISC-V also has an *optional* compressed format for 16b instructions
 - We won't cover that encoding in class. but...
 - It is designed to take "typical" code and produce ~30% more compact code
 - In order to "future proof" things, the 32b instruction encoding needs to assume the possibility of 16b instructions.
 - In the same space you can have 1 32b instruction or 2 16b instructions

Uses of JAL

```
# j pseudo-instruction
j Label = jal x0, Label # Discard return address

# Call function within 218 instructions of PC
jal ra, Label
```

JALR Instruction (I-Format)

31	20 19	15 14	12 11	7 6	0
imm[11:0]	rs1	funct3	rd	opcode	
12 offset[11:0]	5 base	3 0	5 dest	7 JALR	

- JALR rd, rs, immediate
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes

Uses of JALR

```
# ret and jr psuedo-instructions
ret = jr ra = jalr x0, ra, 0
# Call function at any 32-bit absolute address
lui x1, <hi20bits>
jalr ra, x1, <lo12bits>
# Jump PC-relative with 32-bit offset
auipc x1, <hi20bits> # Adds upper immediate value to PC
                        # and places result in x1
jalr x0, x1, <lo12bits>
```

Other uses of JALR

- Pointers to functions...
 - `char (*foo) (char *, char *) = &bar; ...`
`foo(a, b)`
 - Invocation of foo uses a `jalr`
 - Object oriented code (a'la C++/Java)
 - Each object has a pointer to a table of functions
 - To call function `foo()`, look up the right pointer in the table and call it with `jalr`.

Summary of RISC-V Instruction Formats

31	30	25 24	21	20	19	15 14	12 11	8	7	6	0	
												R-type
												I-type
												S-type
												B-type
												U-type
												J-type

Complete RV32I ISA

Computer Science 61C Spring 2018				
imm[31:12]		rd		0110111
imm[31:12]		rd		0010111
imm[20 10:1 11 19:12]		rd		1101111
imm[11:0]	rs1	000	rd	1100111
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]
imm[11:0]	rs1	000	rd	0000011
imm[11:0]	rs1	001	rd	0000011
imm[11:0]	rs1	010	rd	0000011
imm[11:0]	rs1	100	rd	0000011
imm[11:0]	rs1	101	rd	0000011
imm[11:5]	rs2	rs1	000	imm[4:0]
imm[11:5]	rs2	rs1	001	imm[4:0]
imm[11:5]	rs2	rs1	010	imm[4:0]
imm[11:0]	rs1	000	rd	0010011
imm[11:0]	rs1	010	rd	0010011
imm[11:0]	rs1	011	rd	0010011
imm[11:0]	rs1	100	rd	0010011
imm[11:0]	rs1	110	rd	0010011
imm[11:0]	rs1	111	rd	0010011

Wawrynek and Weaver					
0000000	shamt	rs1	001	rd	0010011
0000000	shamt	rs1	101	rd	0010011
0100000	shamt	rs1	101	rd	0010011
0000000	rs2	rs1	000	rd	0110011
0100000	rs2	rs1	000	rd	0110011
0000000	rs2	rs1	001	rd	0110011
0000000	rs2	rs1	010	rd	0110011
0000000	rs2	rs1	011	rd	0110011
0000000	rs2	rs1	100	rd	0110011
0000000	rs2	rs1	101	rd	0110011
0100000	rs2	rs1	101	rd	0110011
0000000	rs2	rs1	110	rd	0110011
0000000	rs2	rs1	111	rd	0110011
0000	pred	succ	00000	000	00000
0000	0000	0000	00000	001	00000
0000000000000		00000	000	00000	1110011
0000000000001		00000	000	00000	1110011
csr		rs1	001	rd	1110011
csr		rs1	010	rd	1110011
csr		rs1	011	rd	1110011
csr		zimm	101	rd	1110011
csr		zimm	110	rd	1110011
csr		zimm	111	rd	1110011

Not in CS61C:
Used for OS & Synchronization

000000000000	00000	000	00000	1110011
000000000001	00000	001	00000	1110011
csr	rs1	001	rd	1110011
csr	rs1	010	rd	1110011
csr	rs1	011	rd	1110011
csr	zimm	101	rd	1110011

Clicker Question:

Project 1 (In Modern American Written English)

- How was Project 1?
 - A. 😊
 - B. 😐
 - C. 😢
 - D. 😰
 - E. 💩
- Project 2 part 1 is now out

Announcements...

- Guerilla Section Tonight (Thursday 2/8)
 - 7-9 pm, 20 Barrows
- Midterm Review Tomorrow (Friday 2/9)
 - 6-9 pm, 145 Dwinelle (probable, check Piazza to confirm)
- No Lecture Tuesday (2/13)
- Midterm Tuesday (2/13)
 - 1 double sided, **handwritten** sheet of notes
 - **We provide** a copy of the green sheet.
 - 7-9 pm, see Piazza for room allocation
 - Conflict: You **must** fill out the form by Friday Noon and you **must** bring a copy of your class schedule to the late exam slot
 - DSP: You **must** have contacted Peiji by email by Friday Noon.

Integer Multiplication (1/3)

- Paper and pencil example (unsigned):

$$\begin{array}{r} \text{Multiplicand} & 1000 & 8 \\ \text{Multiplier} & \underline{\times 1001} & 9 \\ & 1000 & \\ & 0000 & \\ & 0000 & \\ & +1000 & \\ \hline & 01001000 & 72 \end{array}$$

- m bits \times n bits = $m + n$ bit product

Integer Multiplication (2/3)

- In RISC-V, we multiply registers, so:
 - 32-bit value \times 32-bit value = 64-bit value
- Multiplication is ***not*** part of standard RISC-V...
 - Instead it is an ***optional*** extra: The compiler needs to produce a series of shifts and adds if the multiplier isn't present
- Syntax of Multiplication (signed):
 - `mul rd, rs1, rs2`
 - `mulh rd, rs1, rs2`
- Multiplies 32-bit values in those registers and returns either the lower or upper 32b result
 - If you do mulh/mul back to back, the architecture can fuse them
 - Also unsigned versions of the above

Integer Multiplication (3/3)

- Example:
 - in C: **a = b * c;**
 - `int64_t a; int32_t b, c;`
 - Aside, these types are defined in C99, in stdint.h
- in RISC-V:
 - let b be **s2**; let c be **s3**; and let a be **s0** and **s1** (since it may be up to 64 bits)
 - **mulh s1, s2, s3**
mul s0, s2, s3

Integer Division (1/2)

- Paper and pencil example (unsigned):

$$\begin{array}{r} \text{Dividend} \\ \underline{-1001} \quad \text{Quotient} \quad \text{Divisor} \end{array} \quad 1000 \mid 1001010$$

10
101
1010
-1000
10 \text{Remainder}

(or Modulo result)

- $\text{Dividend} = \text{Quotient} \times \text{Divisor} + \text{Remainder}$

Integer Division (2/2)

- Syntax of Division (signed):
 - `div rd, rs1, rs2`
`rem rd, rs1, rs2`
 - Divides 32-bit rs1 by 32-bit rs2, returns the quotient (/) for div, remainder (%) for rem
 - Again, can fuse two adjacent instructions
- Example in C: $a = c / d;$ $b = c \% d;$
- RISC-V:
 - $a \leftrightarrow s0; b \leftrightarrow s1; c \leftrightarrow s2; d \leftrightarrow s3$
 - `div s0, s2, s3`
`rem s1, s2, s3`

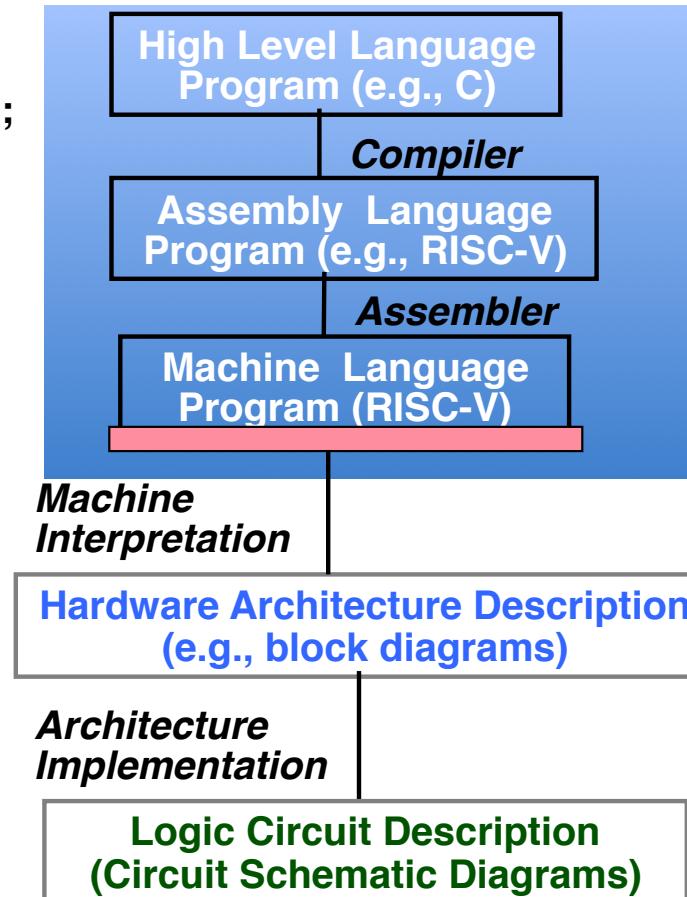
Agenda

- Interpretation vs Compilation
- The CALL chain
- Producing Machine Language

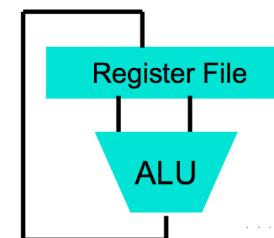
Levels of Representation/Interpretation

```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;  
  
lw t0, 0(a2)  
lw t1, 4(a2)  
sw t1, 0(a2)  
sw t0, 4(a2)
```

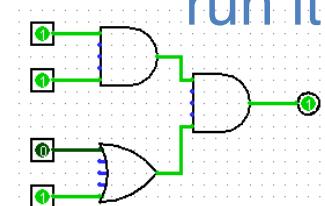
Anything can be represented as a *number*, i.e., data or instructions



0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111

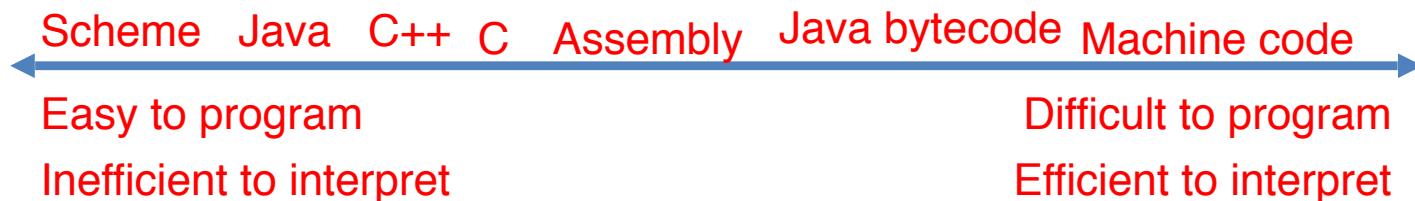


+ How to take a program and run it



Language Execution Continuum

- An **Interpreter** is a program that executes other programs.

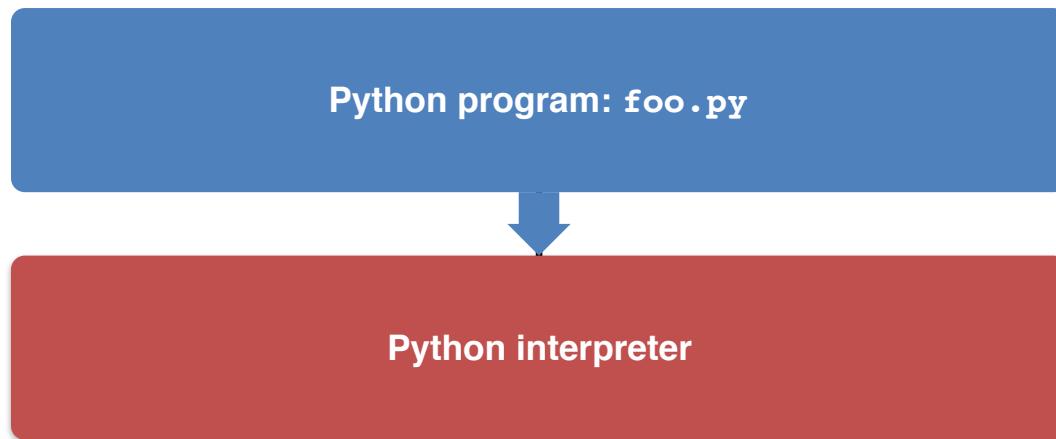


- Language **translation** gives us another option
- In general, we **interpret** a high-level language when efficiency is not critical and **translate** to a lower-level language to increase performance
 - Although this is becoming a “distinction without a difference”
Many interpreters do a “just in time” runtime compilation to bytecode that either is emulated or directly compiled to machine code (e.g. LLVM)

Interpretation vs Translation

- How do we run a program written in a source language?
 - **Interpreter**: Directly executes a program in the source language
 - **Translator**: Converts a program from the source language to an equivalent program in another language
- For example, consider a Python program **foo.py**

Interpretation



- Python interpreter is just a program that reads a python program and performs the functions of that python program
 - Well, that's an exaggeration, the interpreter converts to a simple bytecode that the interpreter runs... Saved copies end up in .pyc files

Interpretation

- Any good reason to interpret machine language in software?
- Simulators: Useful for learning / debugging
- Apple Macintosh conversion
 - Switched from Motorola 680x0 instruction architecture to PowerPC.
 - Similar issue with switch to x86
 - Could require all programs to be re-translated from high level language
 - Instead, let executables contain old and/or new machine code, interpret old code in software if necessary (emulation)

Interpretation vs. Translation? (1/2)

- Generally easier to write interpreter
- Interpreter closer to high-level, so can give better error messages
 - Translator reaction: add extra information to help debugging (line numbers, names):
This is what `gcc -g` does, it tells the compiler to add all the debugging information
- Interpreter slower (10x?), code smaller (2x? or not?)
- Interpreter provides instruction set independence: run on any machine

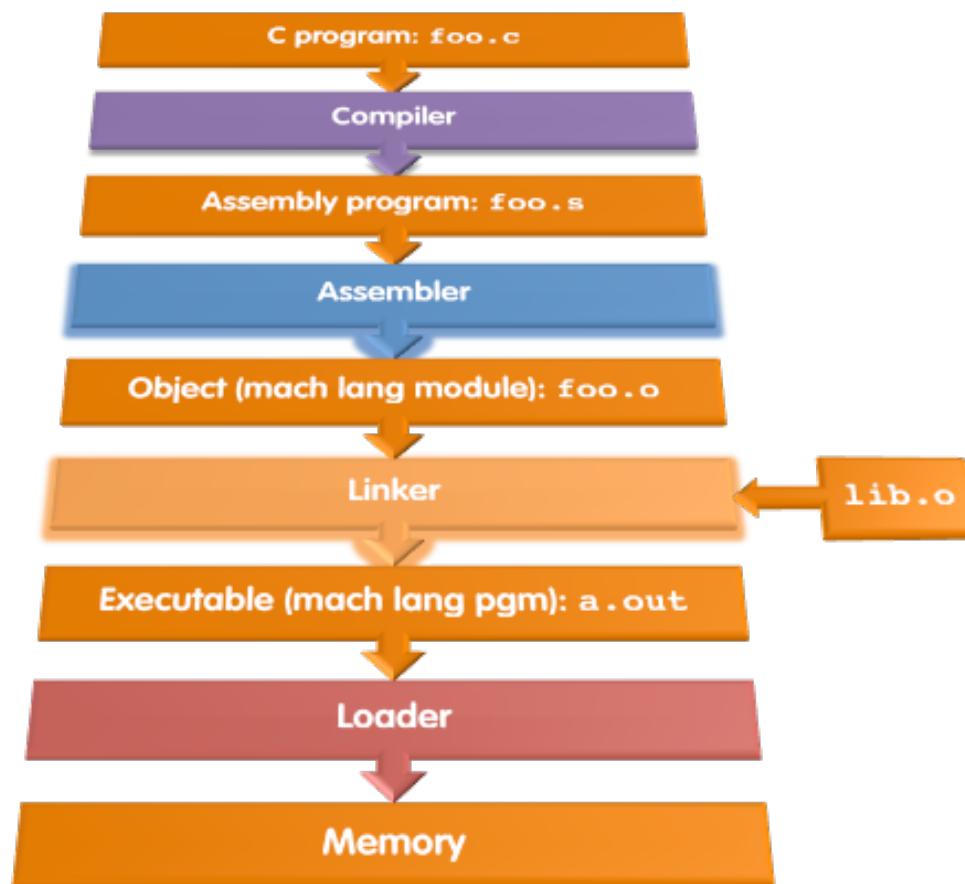
Interpretation vs. Translation? (2/2)

- Translated/compiled code almost always more efficient and therefore higher performance:
 - Important for many applications, particularly operating systems.
 - Compiled code does the hard work once: during compilation
 - Which is why most “interpreters” these days are really “just in time compilers”: don’t throw away the work processing the program

Agenda

- Interpretation vs Compilation
- The CALL chain
- Producing Machine Language

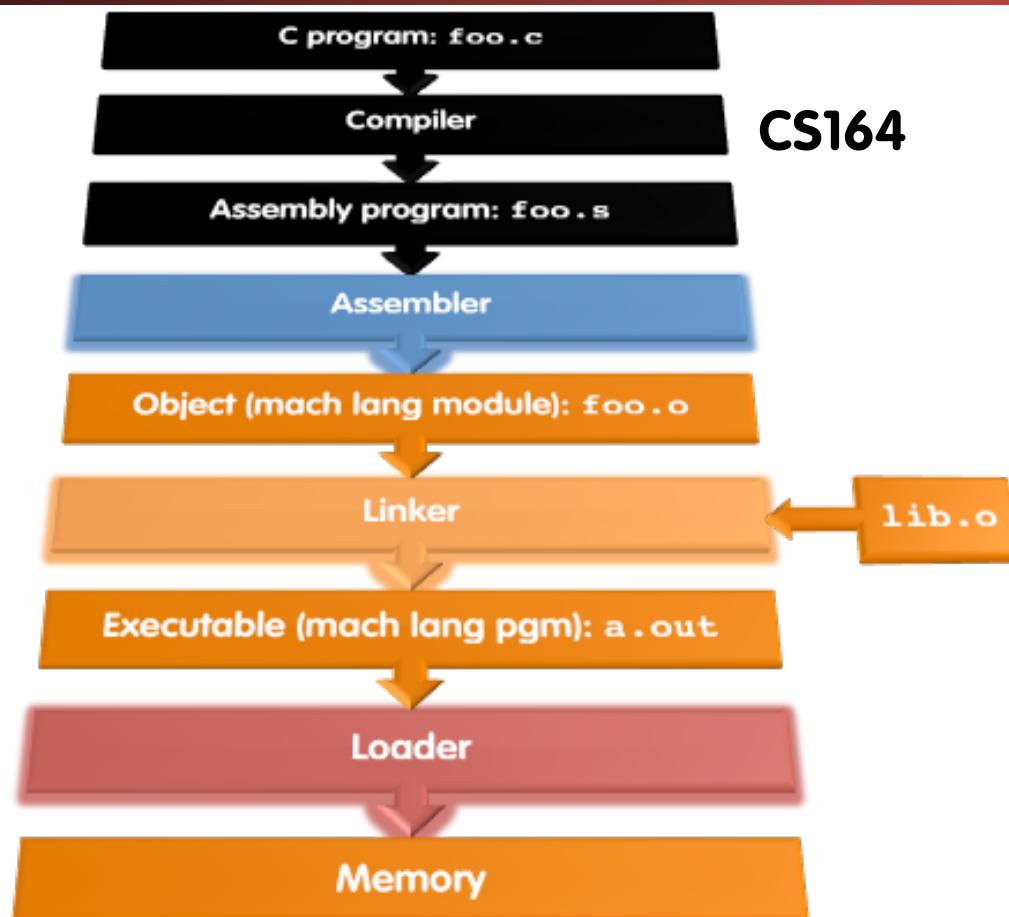
Steps Compiling a C program



Compiler

- Input: High-Level Language Code
(e.g., **foo.c**)
- Output: Assembly Language Code (e.g. MAL)
(e.g., **foo.s** for RISC-V)
 - Code matches the ***calling convention*** for the architecture
- Note: Output *may* contain pseudo-instructions
- Pseudo-instructions: instructions that assembler understands but not in machine
 - **j label** \Rightarrow **jal x0 label**

Where Are We Now?



Assembler

- Input: Assembly Language Code (e.g. MAL)
(e.g., **foo.s**)
- Output: Object Code, information tables (TAL)
(e.g., **foo.o**)
- Reads and Uses **Directives**
- Replace Pseudo-instructions
- Produce Machine Language
- Creates **Object File**

Assembler Directives

- Give directions to assembler, but do not produce machine instructions
 - **.text:** Subsequent items put in user text segment (machine code)
 - **.data:** Subsequent items put in user data segment (binary rep of data in source file)
 - **.globl sym:** declares `sym` global and can be referenced from other files
 - **.string str:** Store the string `str` in memory and null-terminate it
 - **.word w1...wn:** Store the n 32-bit quantities in successive memory words

Pseudo-instruction Replacement

- Assembler treats convenient variations of machine language instructions as if real instructions

Pseudo	Real
<code>nop</code>	<code>addi x0, x0, 0</code>
<code>not rd, rs</code>	<code>xori rd, rs, -1</code>
<code>beqz rs, offset</code>	<code>beq rs, x0, offset</code>
<code>bgt rs, rt, offset</code>	<code>blt rt, rs, offset</code>
<code>j offset</code>	<code>jal x0, offset</code>
<code>ret</code>	<code>jalr x0, x1, offset</code>
<code>call offset</code> (if too big for just a <code>jal</code>)	<code>auipc x6, offset[31:12]</code> <code>jalr x1, x6, offset[11:0]</code>
<code>tail offset</code> (if too far for a <code>j</code>)	<code>auipc x6, offset[31:12]</code> <code>jalr x0, x6, offset[11:0]</code>

So what is "tail" about...

- Often times your code has a convention like this:
 - { ...
 lots of code
 return foo(y);
}
 - It can be a recursive call to `foo()` if this is within `foo()`, or call to a different function...
- So for efficiency...
 - Evaluate the arguments for `foo()` and place them in `a0-a7...`
 - Restore `ra`
 - Pop the stack
 - Then call `foo() with j or tail`
- Then when `foo()` returns, it can return ***directly*** to where it needs to return to
 - Rather than returning to wherever `foo()` was called and returning from there

Agenda

- Interpretation vs Compilation
- The CALL chain
- Producing Machine Language

Producing Machine Language (1/3)

- Simple Case
 - Arithmetic, Logical, Shifts, and so on
 - All necessary info is within the instruction already
- What about Branches?
 - PC-Relative
 - So once pseudo-instructions are replaced by real ones, we know by how many instructions to branch
- So these can be handled

Producing Machine Language (2/3)

- “Forward Reference” problem
 - Branch instructions can refer to labels that are “forward” in the program:

```
    or    s0, x0, x0
L1: slt   t0, x0, $a1
    beq   t0, x0, L2
    addi  a1, a1, -1
    jal   x0, L1
L2: add   $t1, $a0, $a1
```

- Solved by taking 2 passes over the program
 - First pass remembers position of labels
 - Second pass uses label positions to generate code

Producing Machine Language (3/3)

- What about jumps (**j** and **jal**)?
 - Jumps within a file are PC relative (and we can easily compute)
 - Jumps to **other** files we can't
- What about references to static data?
 - **la** gets broken up into **lui** and **addi**
 - These will require the full 32-bit address of the data
 - These can't be determined yet, so we create two tables...

Symbol Table

- List of “items” in this file that may be used by other files
- What are they?
 - Labels: function calling
 - Data: anything in the `.data` section; variables which may be accessed across files

Relocation Table

- List of “items” this file needs the address of later
- What are they?
 - Any external label jumped to: **jal**
 - external (including lib files)
 - Any piece of data in static section
 - such as the **la** instruction

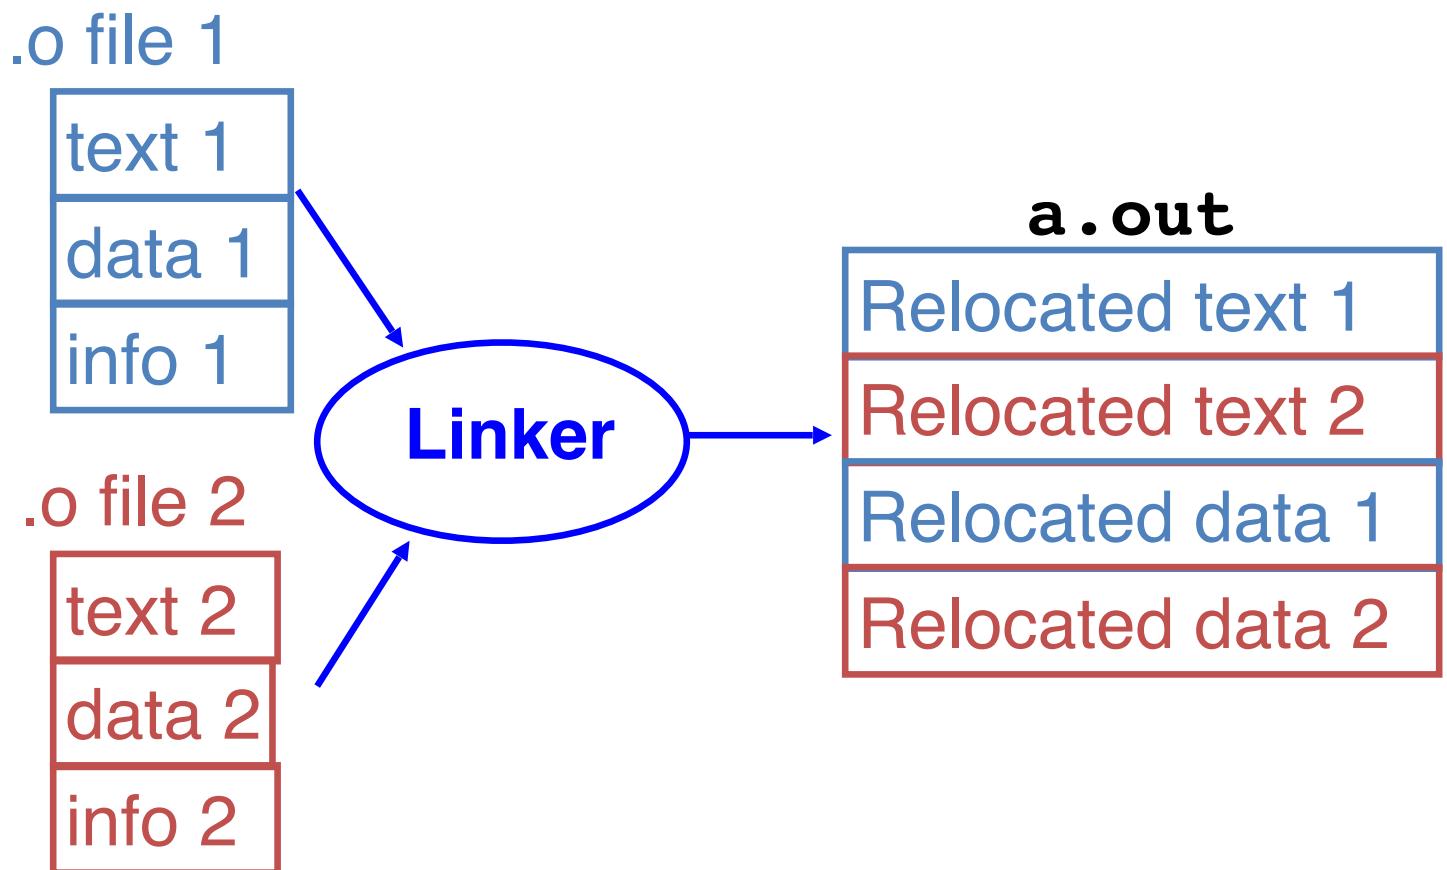
Object File Format

- object file header: size and position of the other pieces of the object file
- text segment: the machine code
- data segment: binary representation of the static data in the source file
- relocation information: identifies lines of code that need to be fixed up later
- symbol table: list of this file's labels and static data that can be referenced
- debugging information
- A standard format is ELF (except Microsoft)
http://www.skyfree.org/linux/references/ELF_Format.pdf

Linker (1/3)

- Input: Object code files, information tables (e.g., `foo.o`, `libc.o`)
- Output: Executable code
(e.g., `a.out`)
- Combines several object (`.o`) files into a single executable (“[linking](#)”)
- Enable separate compilation of files
 - Changes to one file do not require recompilation of the whole program
 - Windows 7 source was > 40 M lines of code!
 - Old name “Link Editor” from editing the “links” in jump and link instructions

Linker (2/3)



Linker (3/3)

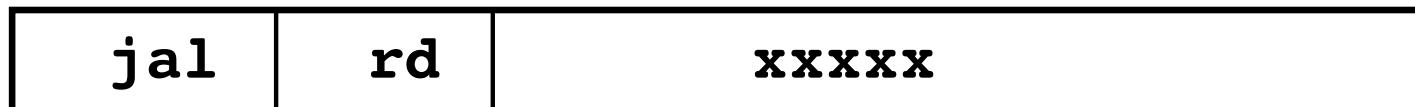
- Step 1: Take text segment from each .o file and put them together
- Step 2: Take data segment from each .o file, put them together, and concatenate this onto end of text segments
- Step 3: Resolve references
 - Go through Relocation Table; handle each entry
 - That is, fill in all **absolute addresses**

Three Types of Addresses

- PC-Relative Addressing (**beq**, **bne**, **jal**)
 - never relocate
- External Function Reference (usually **jal**)
 - always relocate
- Static Data Reference (often **auipc** and **addi**)
 - always relocate
 - RISC-V often uses **auipc** rather than **lui** so that a big block of stuff can be further relocated as long as it is fixed relative to the pc

Absolute Addresses in RISC-V

- Which instructions need relocation editing?
 - Jump and link: ONLY for external jumps



- Loads and stores to variables in static area, relative to the global pointer



- What about conditional branches?



- PC-relative addressing **preserved** even if code moves

Resolving References (1/2)

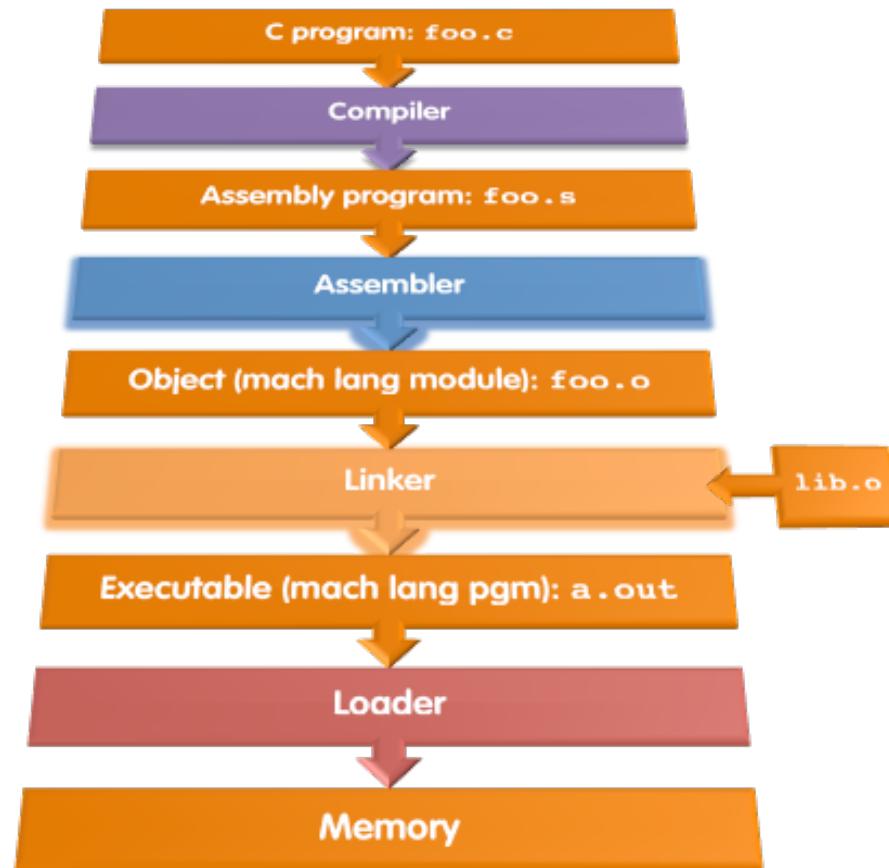
- Linker **assumes** first word of first text segment is at address **0x04000000**.
 - (More later when we study “virtual memory”)
- Linker knows:
 - length of each text and data segment
 - ordering of text and data segments
- Linker calculates:
 - absolute address of each label to be jumped to and each piece of data being referenced

Resolving References (2/2)

- To resolve references:
 - search for reference (data or label) in all “user” symbol tables
 - if not found, search library files
(for example, for `printf`)
 - once absolute address is determined, fill in the machine code appropriately
- Output of linker: executable file containing text and data
(plus header)

In Conclusion...

- Compiler converts a single HLL file into a single assembly language file.
- Assembler removes pseudo-instructions, converts what it can to machine language, and creates a checklist for the linker (relocation table). A .s file becomes a .o file.
 - Does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses.
 - Enables separate compilation, libraries that need not be compiled, and resolves remaining addresses
- Loader loads executable into memory and begins execution.



Loader Basics

- Input: Executable Code
(e.g., **a.out** for MIPS)
- Output: (program is run)
- Executable files are stored on disk
- When one is run, loader's job is to load it into memory and start it running
- In reality, loader is the operating system (OS)
 - loading is one of the OS tasks
 - And these days, the loader actually does a lot of the linking

Loader ... what does it do?

- Reads executable file's header to determine size of text and data segments
- Creates new address space for program large enough to hold text and data segments, along with a stack segment
- Copies instructions and data from executable file into the new address space
- Copies arguments passed to the program onto the stack
- Initializes machine registers
 - Most registers cleared, but stack pointer assigned address of 1st free stack location
- Jumps to start-up routine that copies program's arguments from stack to registers & sets the PC
 - If main routine returns, start-up routine terminates program with the exit system call

Clicker/Peer Instruction

At what point in process are all the machine code bits determined for the following assembly instructions:

1) addu x6, x7, x8

2) jal printf

A: 1) & 2) After compilation

B: 1) After compilation, 2) After assembly

C: 1) After assembly, 2) After linking

D: 1) After compilation, 2) After linking

E: 1) After compilation, 2) After loading

Example: C \Rightarrow Asm \Rightarrow Obj \Rightarrow Exe \Rightarrow Run

C Program Source Code: *prog.c*

```
#include <stdio.h>

int main (int argc, char *argv[ ]) {
    int i, sum = 0;
    for (i = 0; i <= 100; i++)
        sum = sum + i * i;
    printf ("The sum of sq from 0 .. 100 is %d\n", sum);
}
```

“*printf*” lives in “*libc*”

Compilation: MAL:

i = t0, sum = a1

```
.text
.align 2
.globl main
main:
    addi sp, sp, -4
    sw ra, 0(sp)
    mv t0, x0
    mv a1, x0
    li t1, 100
    j check
loop:
    mul t2, t0, t0
    add a1, a1, t2
    addi t0, t0, 1
```

Pseudo-
Instructions?

```
check:
    blt t0, t1 loop:
    la $a0, str
    jal printf
    mv a0, x0
    lw ra, 0(sp)
    addi sp, sp 4
    ret
    .data
    .align 0
str:
    .asciiz "The sum of sq from 0
.. 100 is %d\n"
```

Compilation: MAL:

i = t0, sum = a1

```
.text
.align 2
.globl main
main:
    addi sp, sp, -4
    sw ra, 0(sp)
    mv t0, x0
    mv a1, x0
    li t1, 100
    j check
loop:
    mul t2, t0, t0
    add a1, a1, t2
    addi t0, t0, 1
```

Pseudo-
Instructions?
Underlined

```
check:
    blt t0, t1 loop:
    la $a0, str
    jal printf
    mv a0, x0
    lw ra, 0(sp)
    addi sp, sp 4
    ret
    .data
    .align 0
str:
    .asciiz "The sum of sq from 0
.. 100 is %d\n"
```

Assembly step 1: Remove Pseudo Instructions, assign jumps

```
.text
.align 2
.globl main
main:
    addi sp, sp, -4
    sw ra, 0(sp)
    addi t0, x0, 0
    addi a1, x0, 0
    addi t1, x0, 100
    jal x0, 12
loop:
    mul t2, t0, t0
    add a1, a1, t2
    addi t0, t0, 1
```

Pseudo-
Instructions?
Underlined

```
check:
    blt t0, t1 -16
    lui a0, l.str
    addi a0, a0, r.str
    jal printf
    mv a0, x0
    lw ra, 0(sp)
    addi sp, sp 4
    jalr x0, ra
    .data
    .align 0
str:
    .asciiz "The sum of sq from 0
.. 100 is %d\n"
```

Assembly step 2

Create relocation table and symbol table

- Symbol Table

Label	address (in module)	type
main:	0x00000000	global text
loop:	0x00000014	local text
str:	0x00000000	local data

- Relocation Information

Address	Instr.	type	Dependency
0x0000002c	lui		l.str
0x00000030	addi		r.str
0x00000034	jal		printf

Assembly step 4

- Generate object (.o) file:
 - Output binary representation for
 - text segment (instructions)
 - data segment (data)
 - symbol and relocation tables
 - Using dummy “placeholders” for unresolved absolute and external references
- And next time... We link!