

Input/Output

A Computer is Useless without I/O

- I/O handles persistent storage
 - Disks, SSD memory, etc
- I/O handles user interfaces
 - Keyboard/mouse/display
- I/O handles network

Basic I/O: Devices are Memory

- The memory bus controller ("chipset") doesn't just control the memory...
 - But also controls the PCI I/O devices
- Provides a memory-mapped interface
 - Processor can write & read to control registers to tell the device what to do
 - Processor can write & read data to the device
- The "device driver" runs in the operating system
 - Provides a device independent abstraction to the rest of the OS
 - And user programs have to ask the OS for permission
- Programmed I/O:
 - The CPU is responsible to transfer data to the devices as well as commands

So What Happens When Data Comes?

- I/O is asynchronous
 - It may occur at any time without coordination with what the OS is doing
- Option 1: Trigger an interrupt
 - Jumps control to the interrupt handler which has to figure out what to do...
- Option 2: Wait for the OS to poll the device
 - Do nothing, its the OS's job to check whether something is available

More on Interrupt-Driven

- Highly responsive
 - When data comes in, the interrupt triggers
- Interesting efficiency tradeoff:
 - For low rate its **very** efficient
 - The computer is doing other things except when data comes
 - For high rates its inefficient
 - Interrupts are relatively expensive!
You effectively **have** to nuke all cache state going to/from the OS
- Common design:
 - Interrupt drive by default
 - But if real fast is needed, poll instead
 - Multi-gigabit Network interfaces

Working with real devices

- “Memory mapped I/O”: Device control/data registers mapped to CPU address space
- CPU synchronizes with I/O device:
 - Polling
 - Interrupts
- “Programmed I/O”
 - CPU execs lw/sw instructions for all data movement to/from devices
 - CPU spends time doing 2 things:
 - Getting data from device to main memory
 - Using data to compute

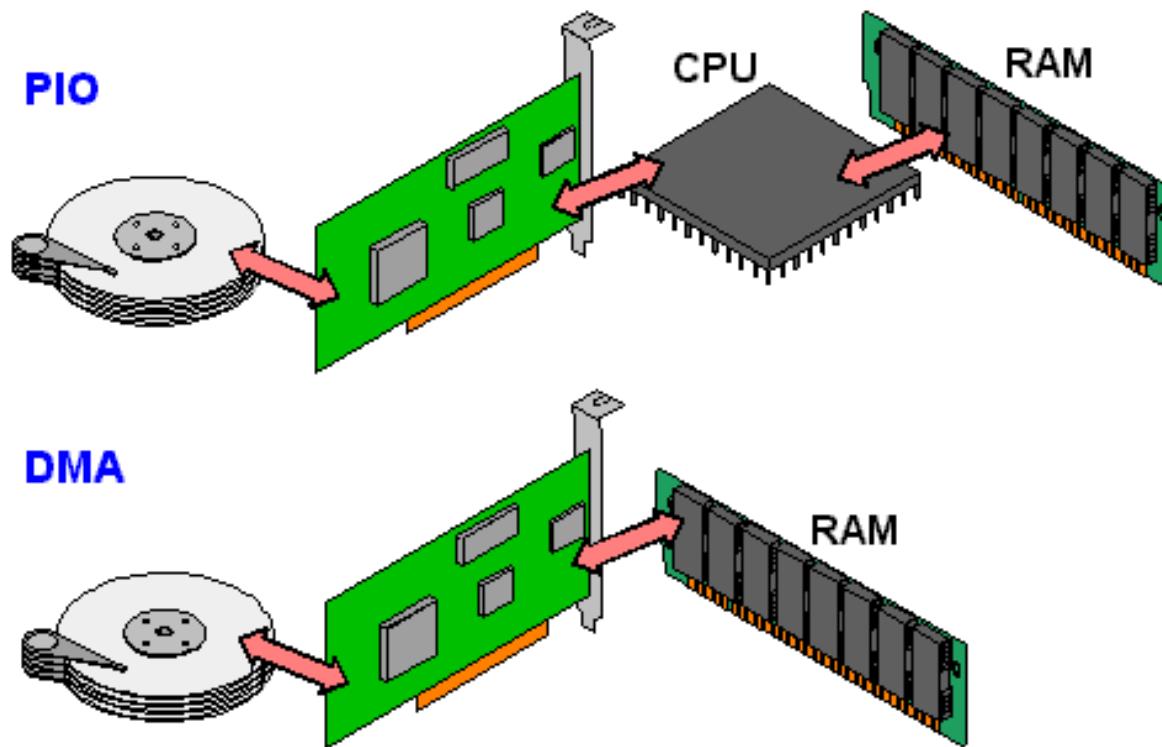
Working with real devices

- “Memory mapped I/O”: Device control/data registers mapped to CPU address space
- CPU synchronizes with I/O device:
 - Polling
 - Interrupts
- “Programmed I/O”: DMA
 - CPU execs lw/sw instructions for all data movement to/from devices
 - CPU spends time doing 2 things:
 - Getting data from device to main memory
 - Using data to compute

What's wrong with Programmed I/O?

- Not ideal because ...
 - CPU has to execute all transfers, could be doing other work
 - Device speeds don't align well with CPU speeds
 - Energy cost of using beefy general-purpose CPU where simpler hardware would suffice
- Until now CPU has sole control of main memory

PIO vs. DMA



Direct Memory Access (DMA)

- Allows I/O devices to directly read/write main memory
- New Hardware: the DMA Engine
- DMA engine contains registers written by CPU:
 - Memory address to place data
 - # of bytes
 - I/O device #, direction of transfer
 - unit of transfer, amount to transfer per burst

Operation of a DMA Transfer

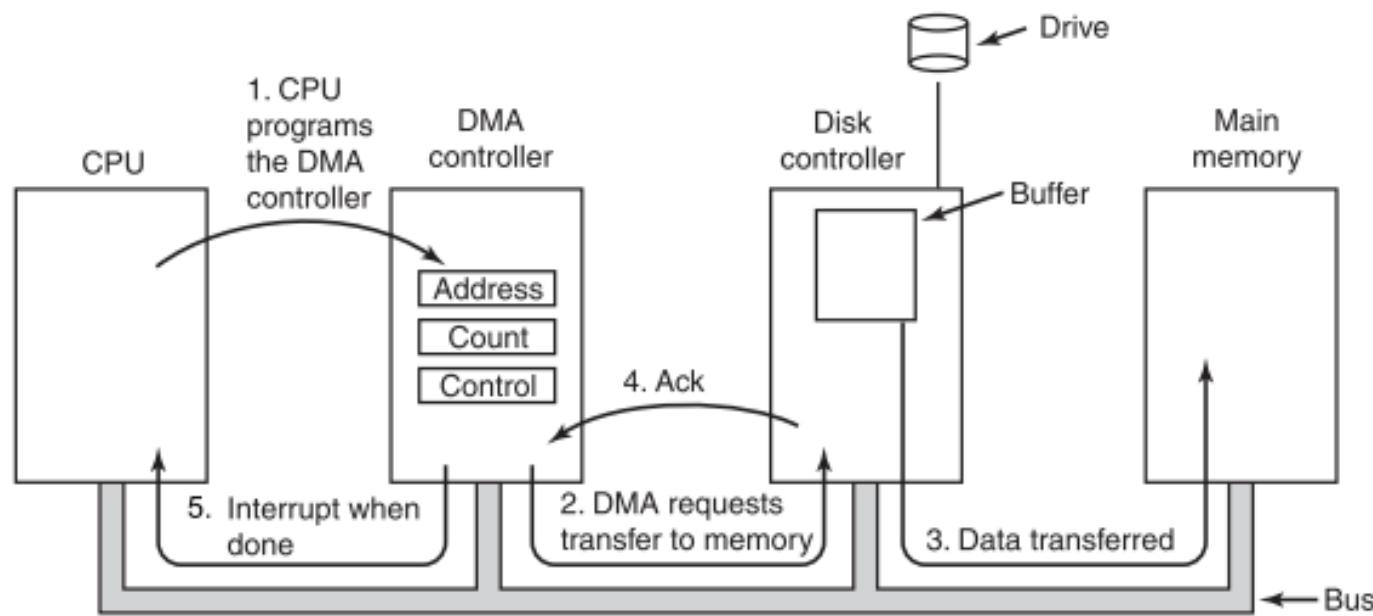


Figure 5-4. Operation of a DMA transfer.

[From Section 5.1.4 Direct Memory Access in *Modern Operating Systems* by Andrew S. Tanenbaum, Herbert Bos, 2014]

DMA: Incoming Data

- Receive interrupt from device
- CPU takes interrupt, begins transfer
 - Instructs DMA engine/device to place data @ certain address
 - Even smarter devices can dispense with this first interrupt:
Agreement with O/S that a particular hunk of memory is writeable by the device,
Device just starts writing where it is supposed to
 - Device/DMA engine handle the transfer
 - CPU is free to execute other things
 - Upon completion, Device/DMA engine interrupt the CPU again

DMA: Outgoing Data

- CPU decides to initiate transfer, confirms that external device is ready
- CPU begins transfer
 - Instructs DMA engine/device that data is available @ certain address
 - Device/DMA engine handle the transfer
 - CPU is free to execute other things
 - Device/DMA engine interrupt the CPU again to signal completion

DMA: Some new problems

- Where in the memory hierarchy do we plug in the DMA engine? Two extremes:
 - Between L1 and CPU:
 - Pro: Free coherency
 - Con: Trash the CPU's working set with transferred data
 - Between Last-level cache and main memory:
 - Pro: Don't mess with caches
 - Con: Need to explicitly manage coherency
- Or just treat like another node in a multiprocessor
 - Cache-coherence is supported by most modern multiprocessors

DMA: Some new problems

- How do we arbitrate between CPU and DMA Engine/Device access to memory? Three options:
 - Burst Mode
 - Start transfer of data block, CPU cannot access memory in the meantime
 - Cycle Stealing Mode
 - DMA engine transfers a byte, releases control, then repeats - interleaves processor/DMA engine accesses
 - Transparent Mode
 - DMA transfer only occurs when CPU is not using the system bus

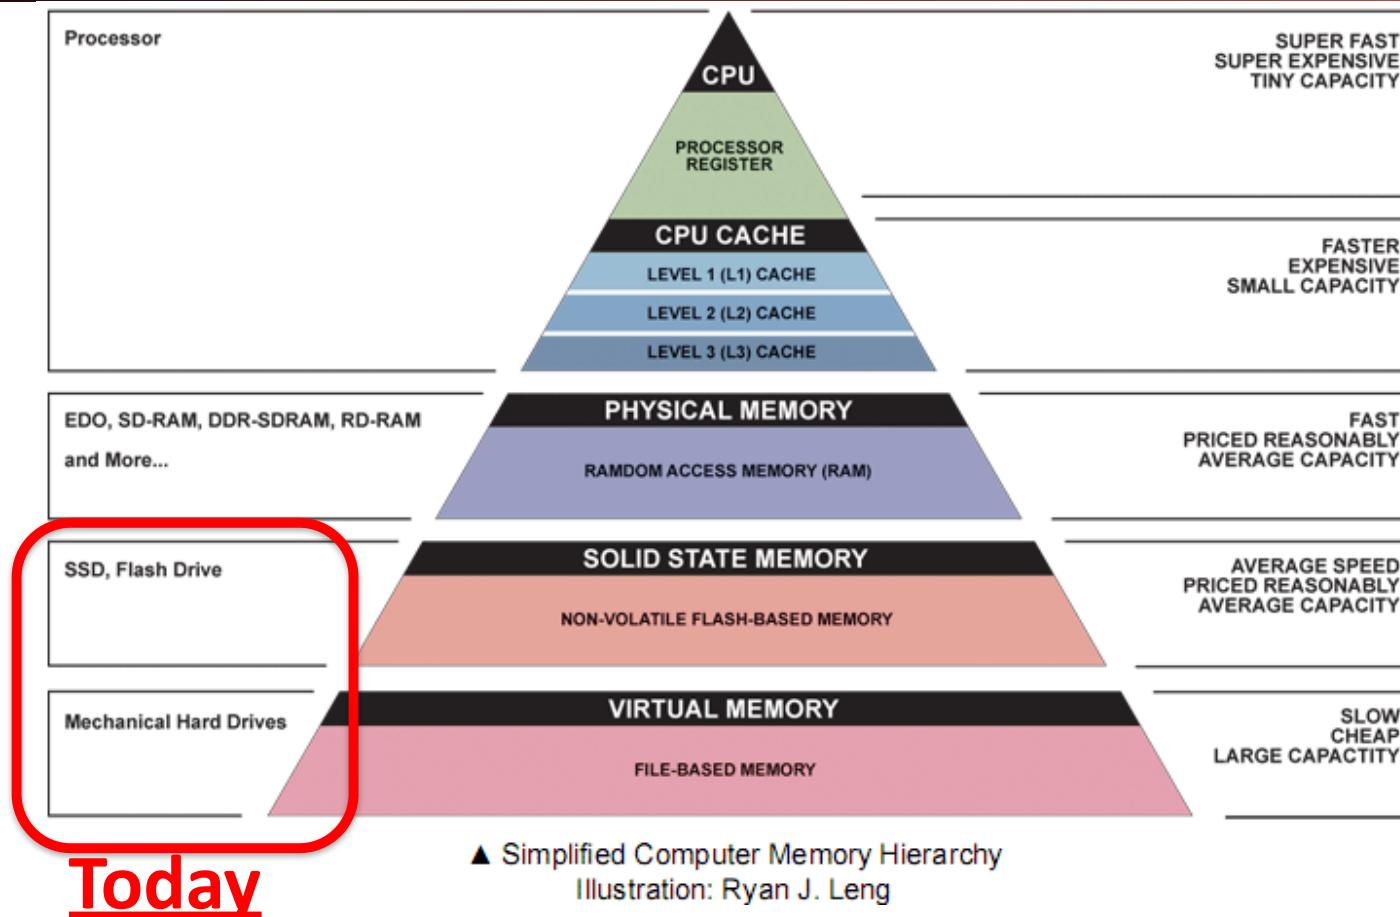
DMA: Cool Interactions with Virtual Memory...

- We want to do high performance I/O for a ***unprivileged process***
 - Normally I/O requires context switching to the Operating System
 - Idea: Map the DMA'ed regions into the user program's memory space!
 - High performance disk: "Memory map" the file. VM says "yeah, the file is in memory here", when you actually want to read it swaps it in from disk using a DMA transfer
 - High performance memory: "User-level buffers". The network card's I/O buffers are mapped into the user process's memory
 - Result is now I/O operations don't even need to ask the OS once things are set up

Announcements

- Project 5-1 is out now
 - 5-2 out soon, both due 4:30 at 23:59.
 - Start **now**: The big problem is making sure you get your environment working.
- HW5 due next Friday
- Guerilla section next thursday

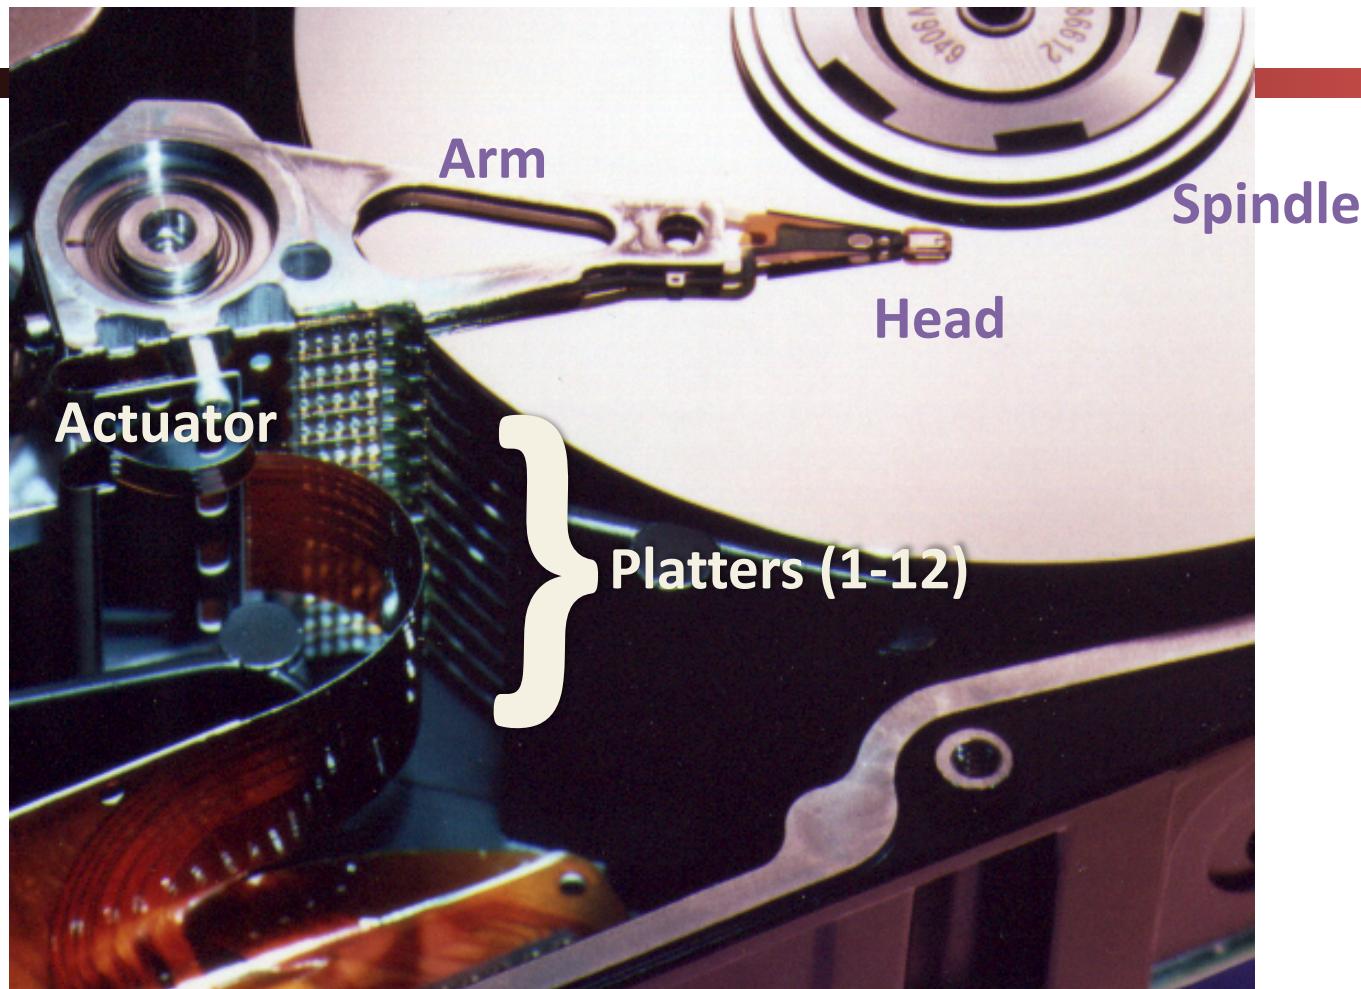
Computer Memory Hierarchy



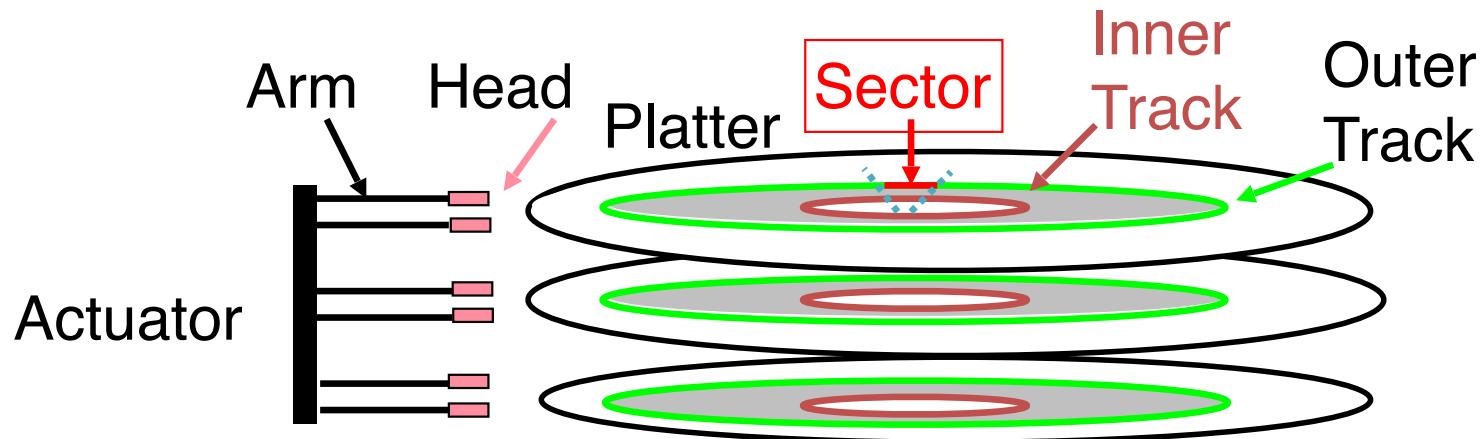
Magnetic Disk – common I/O device

- A kind of computer memory
 - Information stored by magnetizing ferrite material on surface of rotating disk
 - Similar to tape recorder except digital rather than analog data
- A type of non-volatile storage
 - Retains its value without applying power to disk.
- Two Types of Magnetic Disk
 - Hard Disk Drives (HDD) – faster, more dense, non-removable.
 - Floppy disks – slower, less dense, removable (now replaced by USB “flash drive”).
- Purpose in computer systems (Hard Drive):
 - Working file system + long-term backup for files
 - Secondary “backing store” for main-memory. Large, inexpensive, slow level in the memory hierarchy (virtual memory)

Photo of Disk Head, Arm, Actuator



Disk Device Terminology

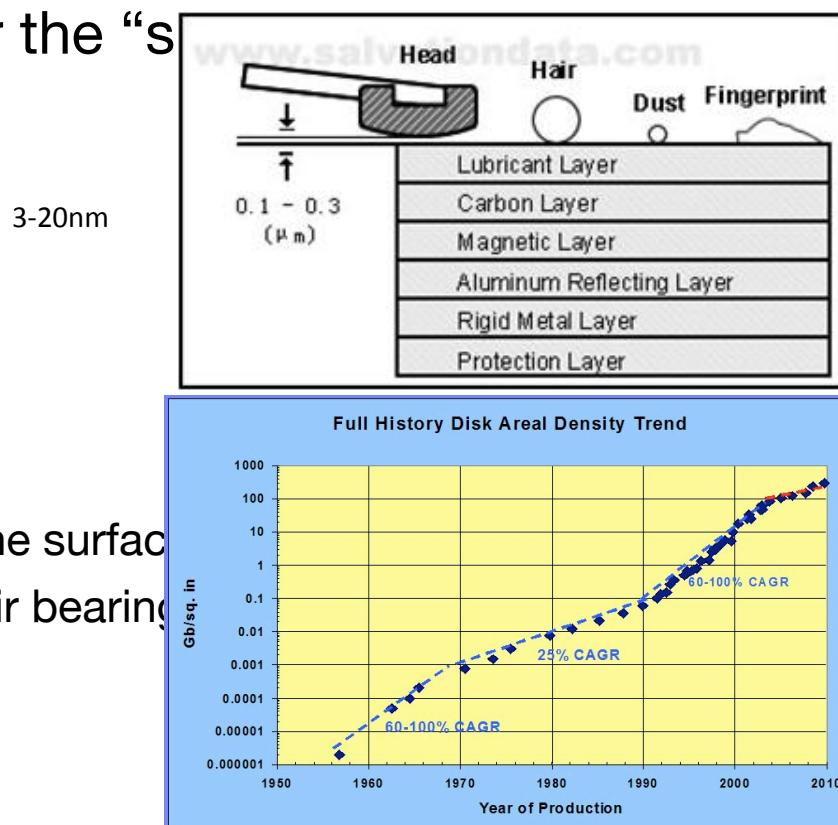


- Several platters, with information recorded magnetically on both surfaces (usually)
- Bits recorded in tracks, which in turn divided into sectors (e.g., 512 Bytes)
- Actuator moves head (end of arm) over track (“seek”), wait for sector rotate under head, then read or write

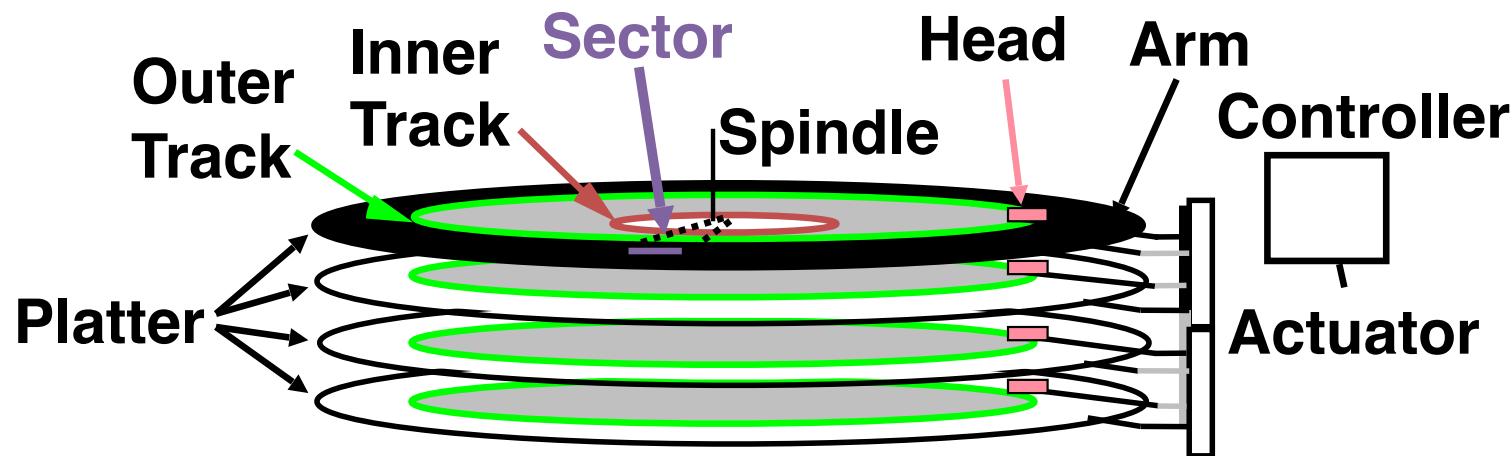
[Video of hard disk in action](#)

Hard Drives are Sealed

- The closer the head to the disk, the smaller the “space” between the head and the disk surface, making the recording denser.
- Measured in Gbit/in²
- ~900 Gbit/in² is state of the art
- Started out at 2 Kbit/in²
- ~450,000,000x improvement in ~60 years
- Disks are sealed to keep the dust out.
 - Heads are designed to “fly” at around 3-20nm above the surface
 - 99.999% of the head/arm weight is supported by the air bearing between the disk and the head
- Some drives are even sealed with Helium
 - Lower drag than air



Disk Device Performance (1/2)



- **Disk Access Time = Seek Time + Rotation Time + Transfer Time + Controller Overhead**
 - Seek Time = time to position the head assembly at the proper cylinder
 - Rotation Time = time for the disk to rotate to the point where the first sectors of the block to access reach the head
 - Transfer Time = time taken by the sectors of the block and any gaps between them to rotate past the head

Disk Device Performance (2/2)

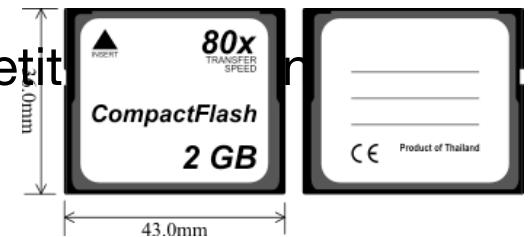
- Average values to plug into the formula:
- Rotation Time: Average distance of sector from head?
 - 1/2 time of a rotation
 - 7200 Revolutions Per Minute \Rightarrow 120 Rev/sec
 - 1 revolution = $1/120$ sec \Rightarrow 8.33 milliseconds
 - 1/2 rotation (revolution) \Rightarrow 4.17 ms
 - Seek time: Average no. tracks to move arm?
 - Number of tracks/3 (see CS186 for the math)
 - Then, seek time = number of tracks moved \times time to move across one track

But wait!

- Performance estimates are different in practice:
- Many disks have on-disk caches, which are completely hidden from the outside world
- Previous formula completely replaced with on-disk cache access time

Where does Flash memory come in?

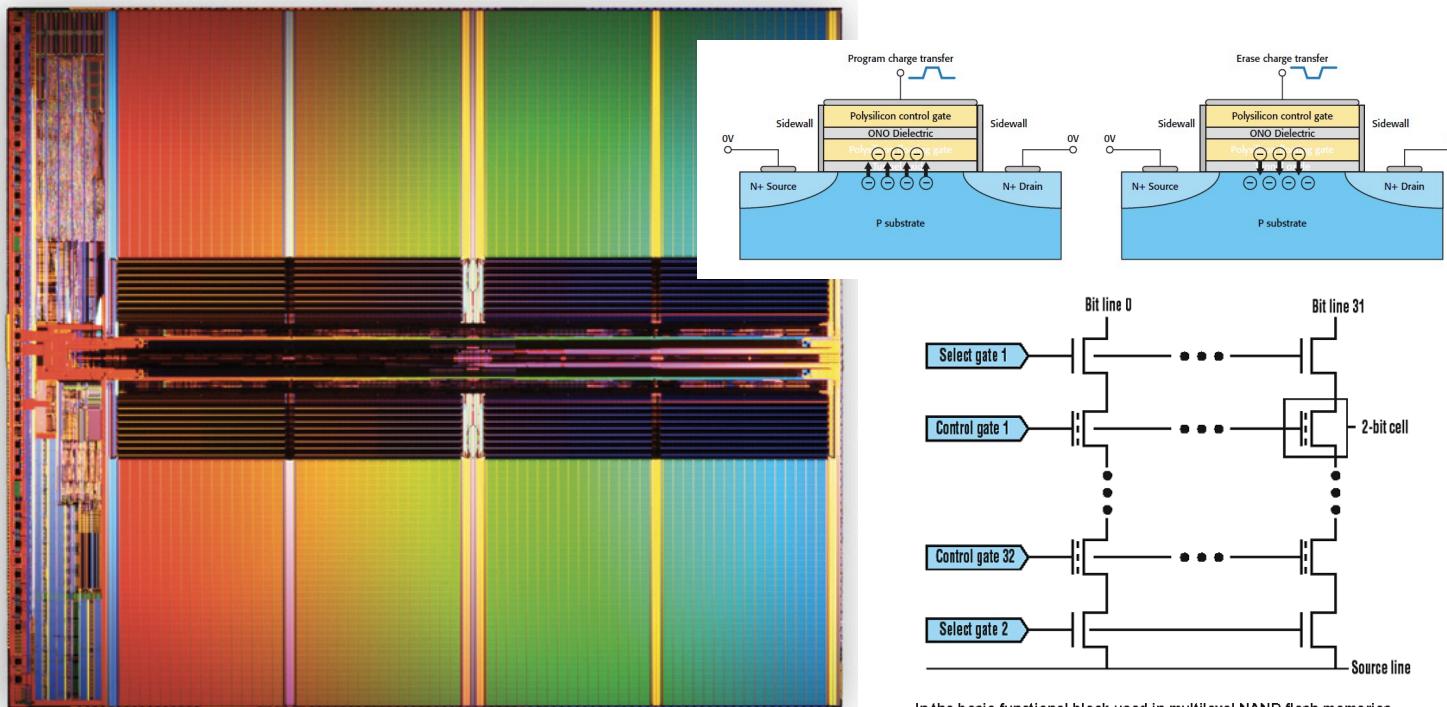
- ~10 years ago: Microdrives and Flash memory (e.g., CompactFlash) went head-to-head
 - Both non-volatile (retains contents without power supply)
 - Flash benefits: lower power, no crashes (no moving parts, need to spin μdrives up/down)
 - Disk cost = fixed cost of motor + arm mechanics, but actual cost very low
 - Flash cost = most cost/bit of flash chips
 - Over time, cost/bit of flash came down, became cost competitive storage



Flash Memory / SSD Technology

Computer Science 61C Spring 2018

Wawrynek and Weaver



2. Micron's triple-level cell (TLC) flash memory stores 3 bits of data in each transistor.

In the basic functional block used in multilevel NAND flash memories, 32 rows of bit lines and 32 control-gate lines form a building block that's repeated many times to form the memory array. The select gate lines are used with the control gate lines to control access to the array.

- NMOS transistor with an additional conductor between gate and source/drain which “traps” electrons. The presence/absence is a 1 or 0
- Memory cells can withstand a limited number of program-erase cycles. Controllers use a technique called *wear leveling* to distribute writes as evenly as possible across all the flash blocks in the SSD.

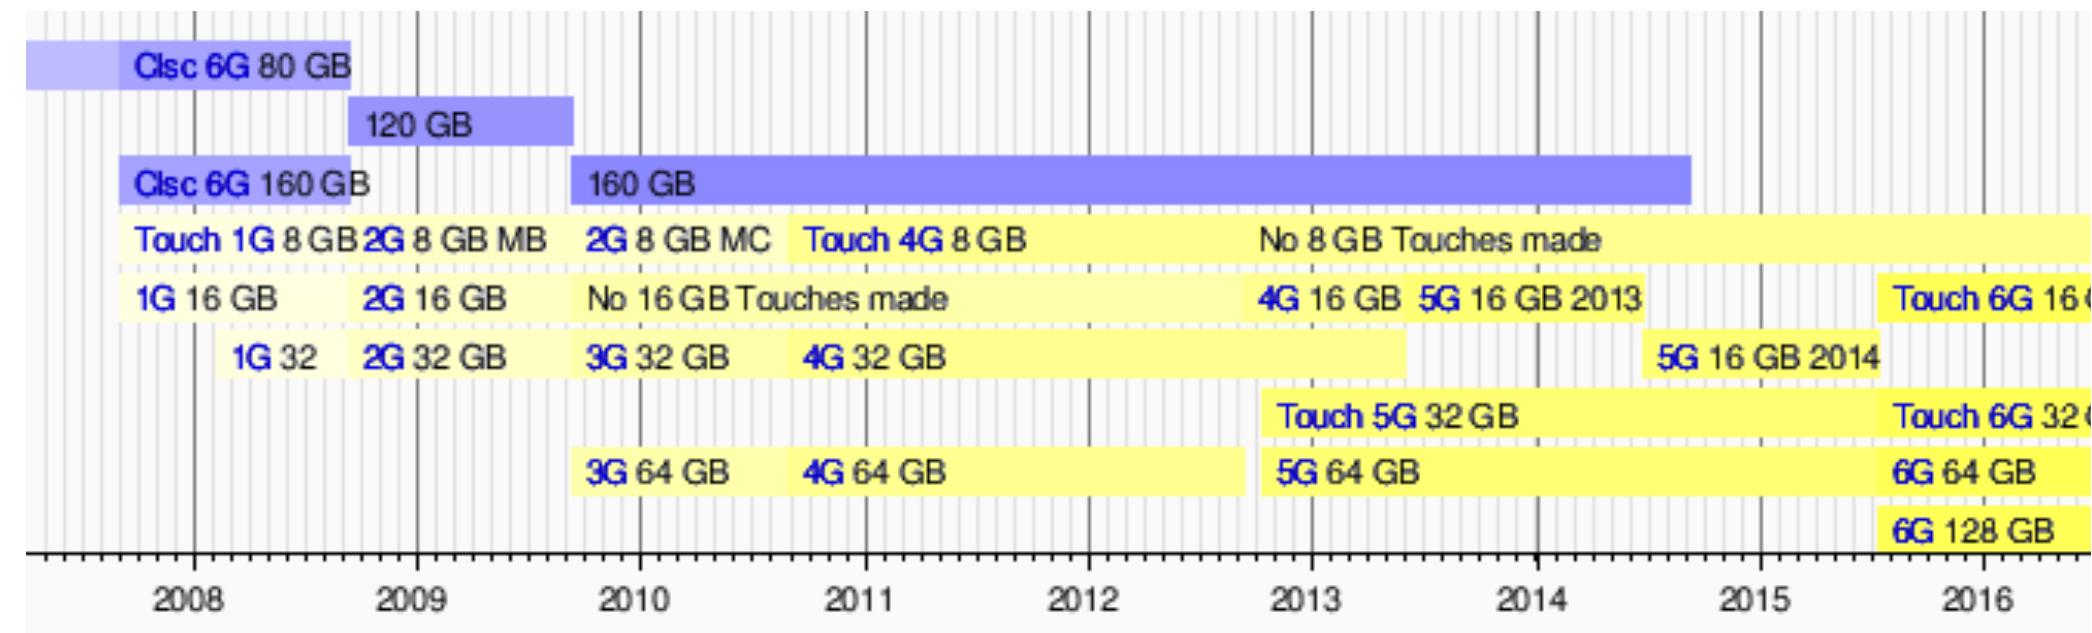
Flash and Latency...

- Flash bandwidth is ~ to spinning disk
 - And spinning disk is still a lot bigger storage/\$ and storage/cc
- But Flash's big party trick: ***no seek time!***
 - No additional latency for random access vs sequential access
- This is huge:
 - Operating system boot ***should be*** as linear as possible
 - Apple spent a huge amount of effort to optimize this on macs...
 - And when I swapped a spinning disk for an SSD, boot times cut in half...

What did Apple put in its iPods?



And How Things Change... iPod Classic v Touch from Wikipedia

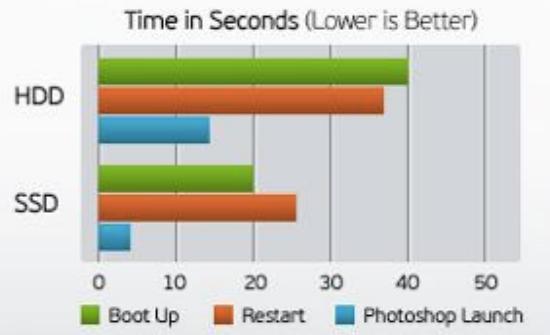


Flash Memory in Smart Phones

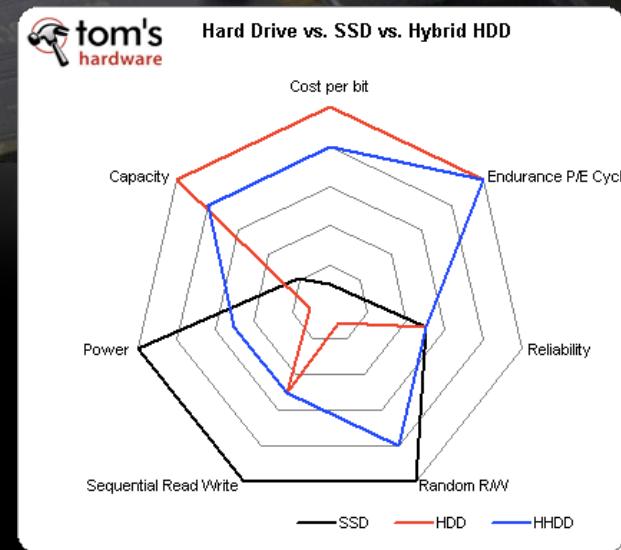
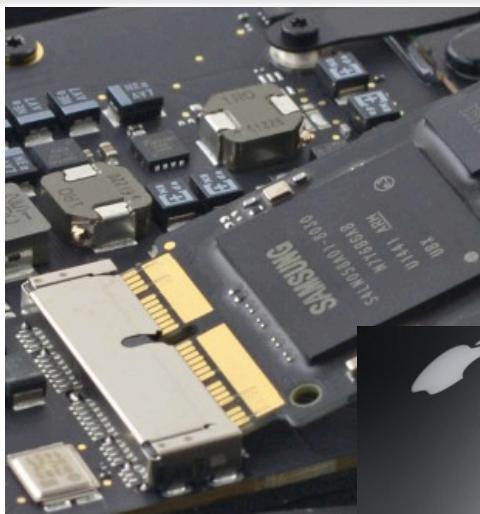
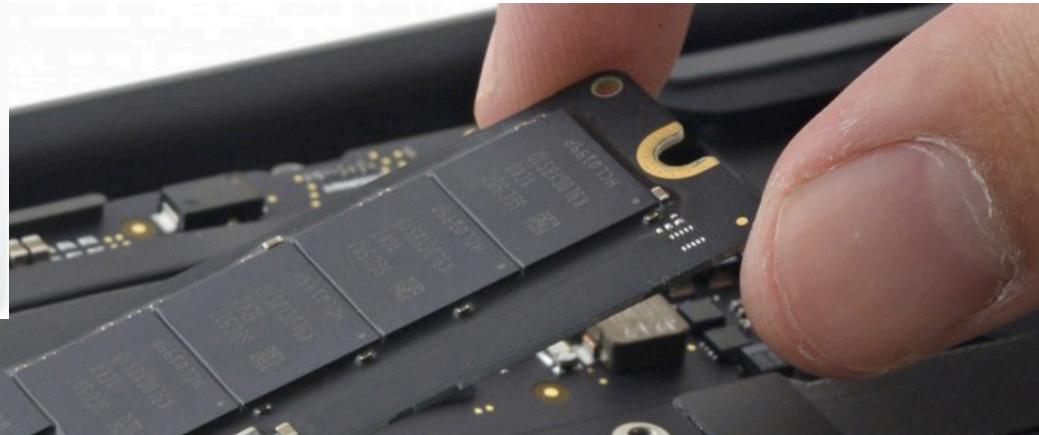


Flash Memory in Laptops – Solid State Drive (SSD)

Computer Science 61C



Wawrynek and Weaver



iClicker Question

- We have the following disk:
 - 15000 Cylinders, 1 ms to cross 1000 Cylinders
 - 15000 RPM = 4 ms per rotation
 - Want to copy 1 MB, transfer rate of 1000 MB/s
 - 1 ms controller processing time
- What is the access time using our model?

Disk Access Time = Seek Time + Rotation Time + Transfer Time + Controller Processing Time

A	B	C	D	E
10.5 ms	9 ms	8.5 ms	11.4 ms	12 ms

iClicker Question

- We have the following disk:

- 15000 Cylinders, 1 ms to cross 1000 Cylinders
- 15000 RPM = 4 ms per rotation
- Want to copy 1 MB, transfer rate of 1000 MB/s
- 1 ms controller processing time

- What is the access time?

Seek = # cylinders/3 * time = $15000/3 * 1\text{ms}/1000 \text{ cylinders} = 5\text{ms}$

Rotation = time for $\frac{1}{2}$ rotation = $4 \text{ ms} / 2 = 2 \text{ ms}$

Transfer = Size / transfer rate = $1 \text{ MB} / (1000 \text{ MB/s}) = 1 \text{ ms}$

Controller = 1 ms

Total = $5 + 2 + 1 + 1 = 9 \text{ ms}$

Networks: Talking to the Outside World

- Originally sharing I/O devices between computers
 - E.g., printers
- Then communicating between computers
 - E.g., file transfer protocol
- Then communicating between people
 - E.g., e-mail
- Then communicating between networks of computers
 - E.g., file sharing, www, ...
- Then turning multiple cheap systems into a single computer
 - Warehouse scale computing

The Internet (1962)

www.computerhistory.org/internet_history

Computer Science 61C Spring 2018

- History
 - 1963: JCR Licklider, while at DoD's ARPA, writes a memo describing desire to connect the computers at various research universities: Stanford, Berkeley, UCLA, ...
 - 1969 : ARPA deploys 4 “nodes” @ UCLA, SRI, Utah, & UCSB
 - 1973 Robert Kahn & Vint Cerf invent TCP, now part of the Internet Protocol Suite
- Internet growth rates
 - Exponential since start
 - But finally starting to hit human scale limits although lots of silicon cockroaches...



www.greatachievements.org/?id=3736
en.wikipedia.org/wiki/Internet_Protocol_Suite

The World Wide Web (1989)

en.wikipedia.org/wiki/History_of_the_World_Wide_Web

Computer Science 61C Spring 2018

Wawrynek and Weaver

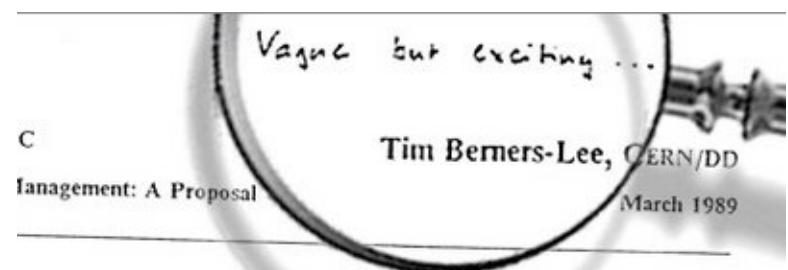
- “System of interlinked hypertext documents on the Internet”
- History
 - 1945: Vannevar Bush describes hypertext system called “memex” in article
 - 1989: Sir Tim Berners-Lee proposed and implemented the first successful communication between a Hypertext Transfer Protocol (HTTP) client and server using the internet.
 - 1993: NCSA Mosaic: A graphical HTTP client
 - ~2000 Dot-com entrepreneurs rushed in, 2001 bubble burst
- Today : Access anywhere!



Tim Berners-Lee



World's First web server in 1990

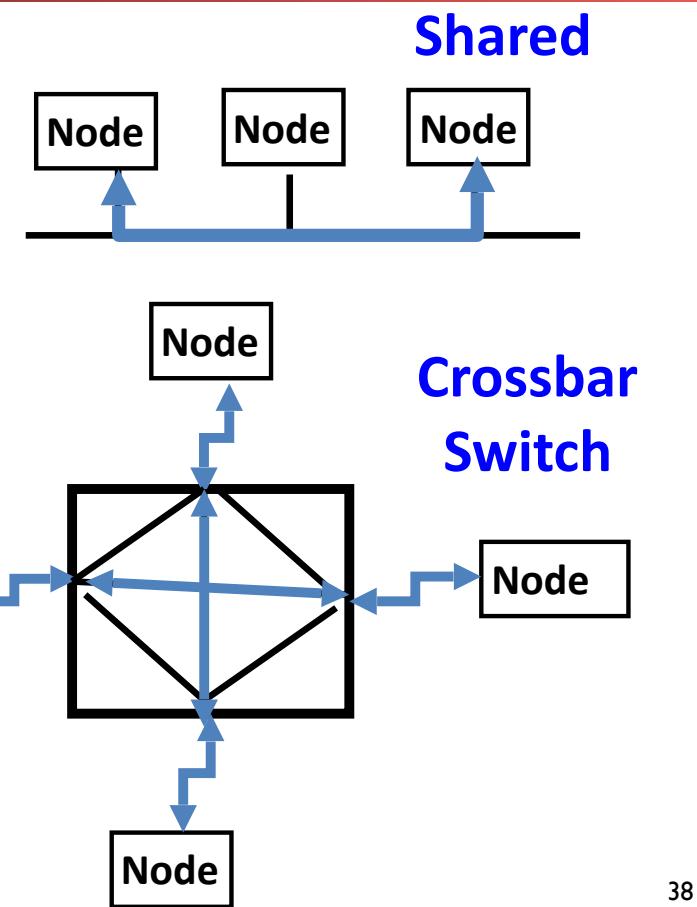


Information Management: A Proposal

Abstract

Shared vs. Switch-Based Networks

- Shared vs. Switched:
 - **Shared:** 1 at a time (CSMA/CD)
 - **Switched:** pairs (“point-to-point” connections) communicate at same time
- Aggregate bandwidth (BW) in switched network is many times that of shared:
 - point-to-point faster since no arbitration, simpler interface
- Wired is almost always switched
- Wireless is by definition shared

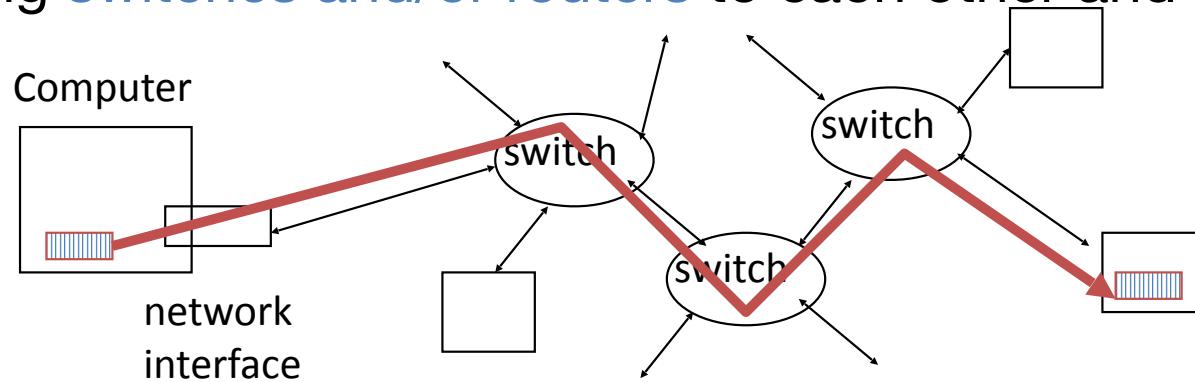


Shared Broadcast

- Old-school Ethernet and Wireless
 - It doesn't just share but all others can see the request?
- How to handle access:
 - Old when I was old skool: Token ring
 - A single "token" that is passed around
 - Ethernet:
 - Listen and send
 - Randomized retry when someone else interrupts you
 - Cable Modem:
 - "Request to send": small request with a listen and send model
 - Big transfers then arbitrated

What makes networks work?

- links connecting switches and/or routers to each other and to computers or devices

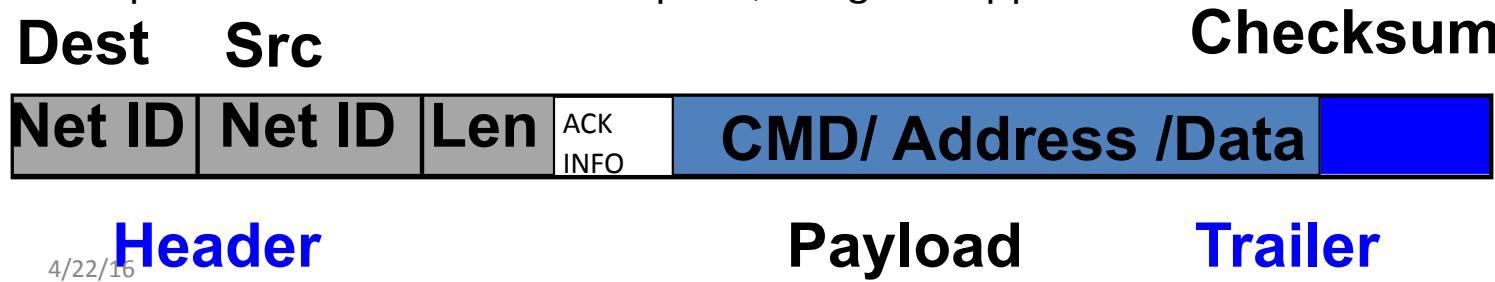


- ability to name the components and to route packets of information - messages - from a source to a destination
- Layering, redundancy, protocols, and encapsulation as means of abstraction (61C big idea)



Software Protocol to Send and Receive

- SW Send steps
 - 1: Application copies data to OS buffer
 - 2: OS calculates checksum, starts timer
 - 3: OS sends DMA request to network interface HW and says start
- SW Receive steps
 - 3: Network interface copies data from network interface HW to OS buffer, triggers interrupt
 - 2: OS calculates checksum, if OK, send ACK; if not, delete message (sender resends when timer expires)
 - 1: If OK, OS copies data to user address space, & signals application to continue



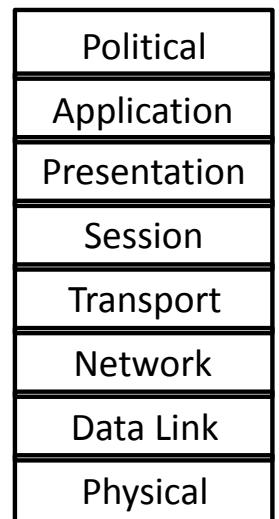
Protocols for Networks of Networks?

What does it take to send packets across the globe?

- Bits on wire or air
- Packets on wire or air
- Delivery packets within a single physical network
- Deliver packets across multiple networks
- Ensure the destination received the data
- Create data at the sender and make use of the data at the receiver

The OSI 7 Layer Network Model

- A conceptual "Stack"
 - Physical Link: eg, the wires/wireless
 - Data Link: Ethernet
 - Network Layer: IP
 - Transport Layer: TCP/UDP
 - Session Layer/Presentation Layer/Application Layer
 - Political Layer: "Feinstein/Burr 'thou shalt not encrypt'"
 - Nick is starting to spend way too much time on "layer 8" issues



Protocol Family Concept

- Protocol: packet structure and control commands to manage communication
- Protocol families (suites): a set of cooperating protocols that implement the network stack
- Key to protocol families is that communication occurs logically at the same level of the protocol, called peer-to-peer...
...but is implemented via services at the next lower level
- Encapsulation: carry higher level information within lower level “envelope”

Inspiration...

- CEO A writes letter to CEO B
 - Folds letter and hands it to assistant
- Assistant:
 - Puts letter in envelope with CEO B's full name
 - Takes to FedEx
- FedEx Office
 - Puts letter in larger envelope
 - Puts name and street address on FedEx envelope
 - Puts package on FedEx delivery truck
- FedEx delivers to other company

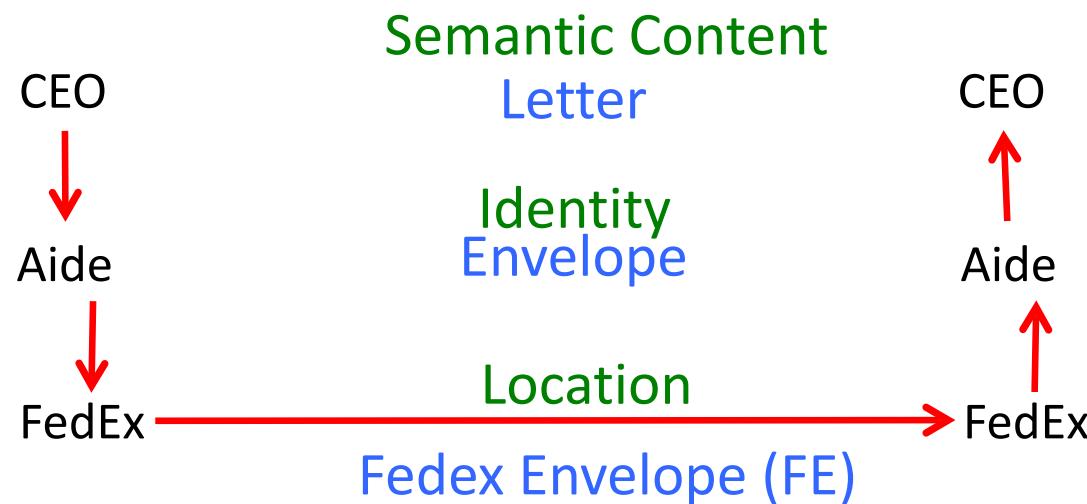
Dear John,

Your days are numbered.

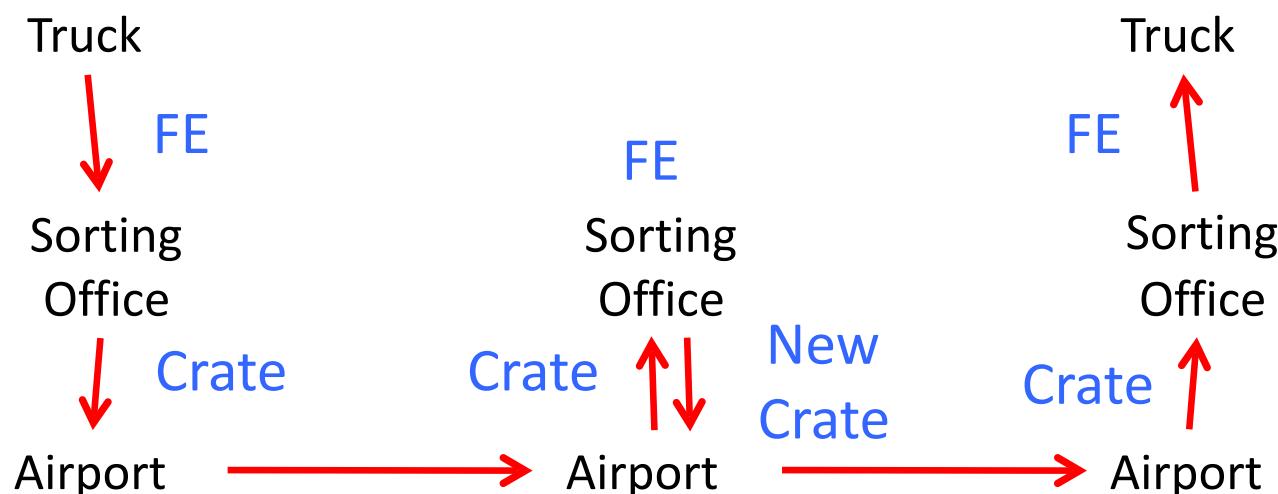
--Pat

The Path of the Letter

“Peers” on each side understand the same things
No one else needs to
Lowest level has most packaging

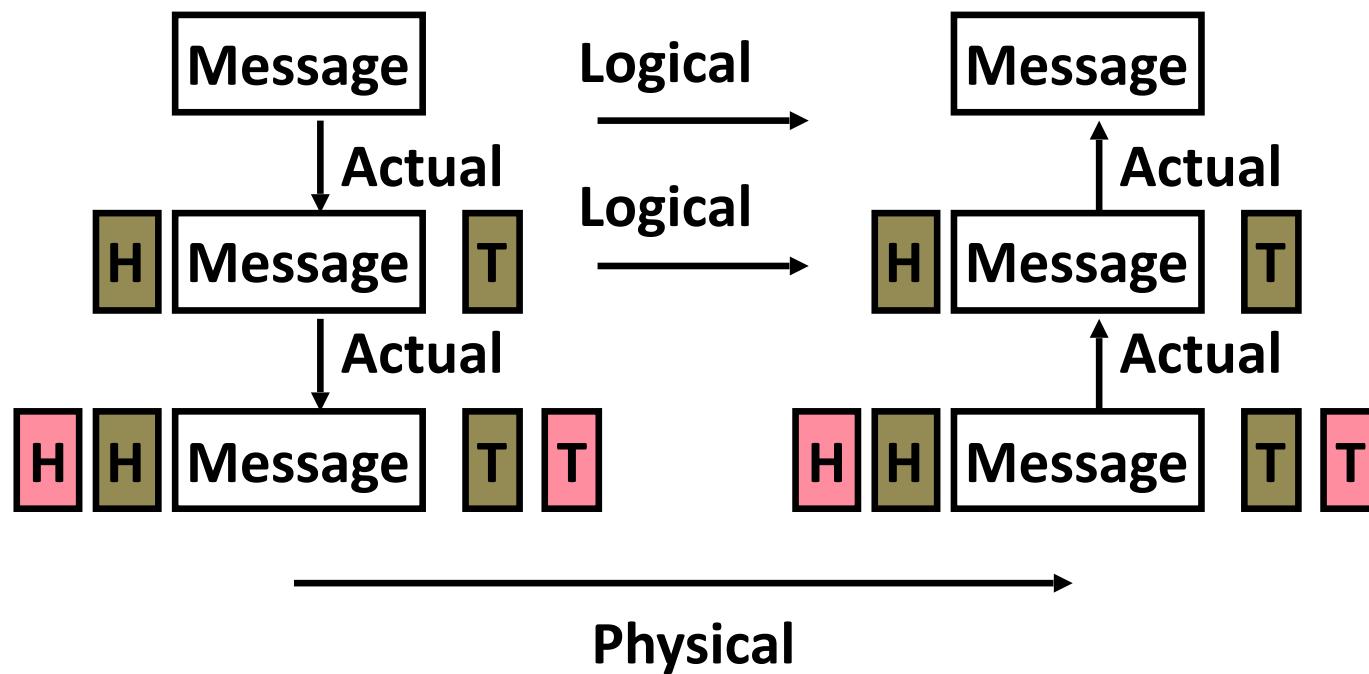


The Path Through FedEx



Deepest Packaging (Envelope+FE+Crate)
at the Lowest Level of Transport

Protocol Family Concept



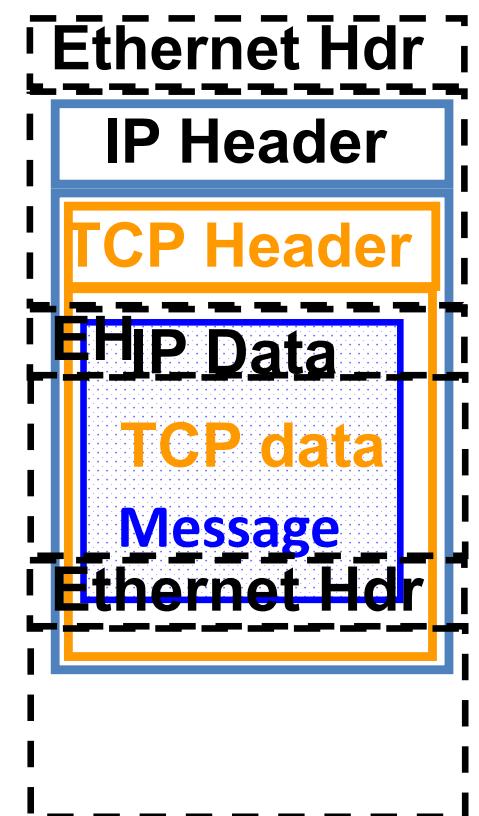
Each lower level of stack “encapsulates” information from layer above by adding header and trailer.

Most Popular Protocol for Network of Networks

- Transmission Control Protocol/Internet Protocol (TCP/IP)
- This protocol family is the basis of the Internet, a WAN (wide area network) protocol
 - IP makes best effort to deliver
 - Packets can be lost, corrupted
 - But corrupted packets should be turned into lost packets
 - TCP guarantees ***reliable, in-order*** delivery
 - TCP/IP so popular it is used even when communicating locally: even across homogeneous LAN (local area network)

TCP/IP packet, Ethernet packet, protocols

- Application sends message
 - TCP breaks into 64KiB segments, adds 20B header
 - IP adds 20B header, sends to network
 - If Ethernet, broken into 1500B packets with headers, trailers



TCP and UDP

The Two Internet Transfer Protocols

- TCP: Transmission Control Protocol
 - Connection based
 - SYN->SYN/ACK->ACK 3-way handshake
 - In order & reliable
 - All data is acknowledged
 - Programming interface is streams of data
- UDP: Universal Datagram Protocol
 - Datagram based
 - Just send messages
 - Out of order & unreliable
 - Datagrams can arrive in any order or be dropped (but not corrupted)
 - Needed for realtime protocols

And Switching Gears: GPIO

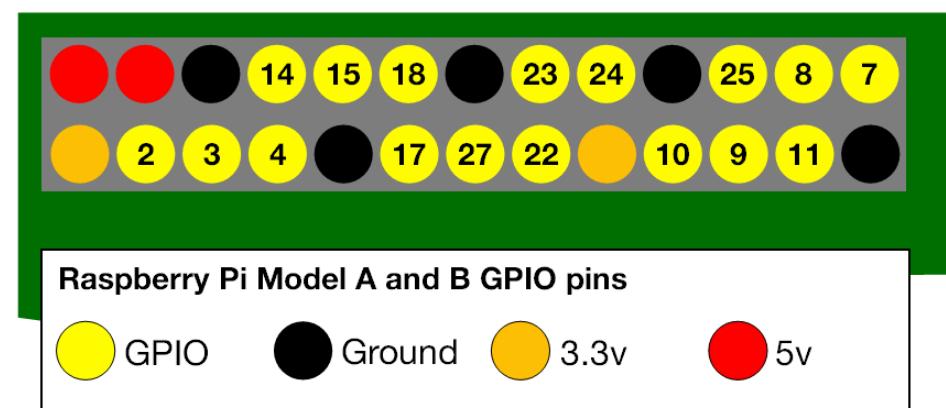
- We see how to do high performance I/O
 - CPU has data it wants to send in main memory
 - Configures device & DMA controller to initiate transfer
 - Device then receives the data through DMA
- We have moderate bandwidth, flexible I/O
 - Universal Serial Bus is really a lightweight network for your slower peripheral devices
- But what about human scale?
 - With people, we only need to react in milliseconds to hours

Reminder: Amdahl's Law and Programming Effort

- Don't optimize where you don't need to
 - And if I only need to react at kHz granularity...
But my processor is a GHz...
- I have 1 million clock cycles to actually decide what to do!
- So lets provide a simple interface

Raspberry Pi GPIO

- A set of physical pins hooked up to the CPU
 - The CPU can write and read these pins as memory, like any other I/O device
- But that is a low level pain for us humans...
 - So the Linux installation provides "files" that can access GPIO
 - You can literally write a 1 or a 0 to a pin or read the value at a pin
- Plus faster & still simple APIs



Using GPIO

- There are a lot of add-on cards...
 - EG, ones for controlling servos
- Or you can build your own
- Combined with USB provides very powerful glue...
- Similarly some even smaller devices:
 - Adafruit "Trinket": 8 MHz 8-bit microcontroller, 5 GPIO pins



“And in conclusion...”

- I/O gives computers their 5 senses
- I/O speed range is 100-million to one
- Polling vs. Interrupts
- DMA to avoid wasting CPU time on data transfers
- Disks for persistent storage, replaced by flash
- Network for communicating to the rest of the planet